

DiceForge Documentation

Mathematics Club, IIT Madras

Pradyumnan Raghuveeran AE22B009

Kulkarni Chinmay Ashish EE22B158

Kailash Gopal D CS22B098

Advik Kabra CS23B004

S Kevin Kinsey EP23B027

Parijat Mukherjee CS23B049

Madhav Bhardwaj CS23B035

Navinkumar L EE23B047

Achintya Raghavan EE23B189

Astha Chand EE23B013

D V Anantha Padmanabh ME23B012

Aditya Sawant CS23B003

Aditi Vaidya ME23B248

Shanddhoshkumar Senthilkumar EE23B116

Contents

1	Introduction	3
2	Functions of the RNGs	4
2.1	Bookkeeping functions	4
2.2	Functions to generate integers	4
2.3	Functions to generate floats	4
2.4	Functions acting on sequences	4
3	Functions of the Distributions	5
3.1	Continuous	5
3.2	Discrete	5
4	Functions for Fitting Data	7
4.1	Fitting to a Cauchy distribution	7
4.2	Fitting to an Exponential distribution	7
4.3	Fitting to an Gaussian distribution	7
4.4	Fitting to a Maxwell distribution	8
4.5	Fitting to a Weibull distribution	8
5	Generators supported	9
5.1	Linear Feedback Shift Register	9
5.1.1	Algorithm	9
5.2	Blum Blum Shub	9
5.2.1	Algorithm	10
5.3	Mersenne Twister	10
5.3.1	Algorithm	10
5.4	XOR Shift	11
5.4.1	Algorithm	11
5.5	Naor-Reingold Pseudo-random Function	12
5.5.1	Algorithm	12
6	Continuous distributions supported	13
6.1	Cauchy	13
6.1.1	Description	13
6.1.2	Parameters	13
6.1.3	Properties	13
6.1.4	Usage	13
6.2	Exponential	13
6.2.1	Description	13
6.2.2	Parameters	14
6.2.3	Properties	14
6.2.4	Usage	14
6.3	Gaussian (Normal)	14
6.3.1	Description	14
6.3.2	Parameters	14
6.3.3	Properties	14
6.3.4	Usage	15
6.4	Maxwell-Boltzmann	15
6.4.1	Description	15

6.4.2	Parameters	15
6.4.3	Properties	15
6.4.4	Usage	15
6.5	Weibull	15
6.5.1	Description	15
6.5.2	Parameters	16
6.5.3	Properties	16
6.5.4	Usage	16
6.6	Custom Distribution	16
6.6.1	Description	16
6.6.2	Parameters	16
6.6.3	Properties	17
6.6.4	Usage	17
7	Discrete distributions supported	18
7.1	Bernoulli	18
7.1.1	Description	18
7.1.2	Parameters	18
7.1.3	Usage	18
7.2	Binomial	18
7.2.1	Description	18
7.2.2	Parameters	18
7.2.3	Usage	18
7.3	Geometric	19
7.3.1	Definition	19
7.3.2	Parameters and PMF	19
7.3.3	Properties	19
7.3.4	Usage	19
7.4	Gibbs	19
7.4.1	Definition and applications	19
7.4.2	Parameters	20
7.4.3	Usage	20
7.5	Hypergeometric	20
7.5.1	Description	20
7.5.2	Parameters	20
7.5.3	Properties	20
7.5.4	Usage	21
7.6	Negative Hypergeometric	21
7.6.1	Description	21
7.6.2	Properties	21
7.6.3	Parameters	21
7.6.4	Usage	21
7.7	Poisson	22
7.7.1	Definition	22
7.7.2	Implementation	22
7.7.3	Parameters	22
7.7.4	Usage	22
8	2D Random Variables	23
8.1	Integrate (1D) Function Usage	23
8.2	Integrate (2D) Function Usage	23

9	Performance and Statistics	24
9.1	Uniformity of the PRNGs	24
9.2	Time Performance	24
9.3	Dieharder Tests Results	25
9.3.1	Mersenne Twister	25
9.3.2	XORShift	27
9.3.3	LFSR	29
9.3.4	Naor Reingold (32 bit only)	31
9.3.5	Blum Blum Shub	32
10	Examples	34
10.1	Fitting data to a Cauchy distribution	34
10.2	Estimating the value of π using Monte-Carlo methods	34
10.3	Path Tracer using Mersenne Twister to sample BRDFs	35
10.4	Chaos Game using LFSR to choose points	35
10.5	Sampling a Maxwell-Boltzmann distribution	35
11	Appendix	36
11.1	Parameters in Mersenne Twister	36
12	Sources	37

List of Figures

1	Fitting to a distribution	34
2	Monte-Carlo estimation of π	34
3	A Path-Tracer using DiceForge's Mersenne Twister to sample light-ray directions	35
4	A Sierpinski Triangle formed through a Chaos Game using a regular triangle and a factor of $\frac{1}{2}$	35
5	PDF from sampled Maxwell-Boltzmann	36

List of Tables

1	Uniformity of the Generators	24
2	Time taken to generate 10^8 random numbers by DiceForge	24
3	Time taken to generate 10^8 random numbers by other libraries	24

1 Introduction

DiceForge is a C++ library with the same functionality as Python's random library with more features useful in scientific computing.

DiceForge's functionality can be split into three major sections - pseudo random number generation, sampling probability distributions, and fitting data to standard probability distributions.

Blum Blum Shub, Linear Feedback Shift Register, Mersenne Twister, XOR Shift, and Naor-Reingold pseudo-random function are the supported algorithms for generating pseudo random numbers. All the algorithms can generate:

- (a) A uniformly selected non-zero 64-bit or 32-bit unsigned integer.
- (b) A uniformly selected floating point number between 0 and 1
- (c) A uniformly selected integer between two integers a and b (both inclusive).
- (d) A uniformly selected float floating point number between two values a and b.

In addition to the uniform distribution (which is an inbuilt function of the generators themselves), DiceForge supports the generation of samples from Cauchy, Exponential, Gaussian (Normal), Maxwell, Weibull, Bernoulli, Binomial, Geometric, Gibbs, Hypergeometric, Negative-Hypergeometric, and Poisson distributions. A class 'CustomDistribution' is also featured for sampling random variables from user defined distributions for any given valid PDF.

Currently, DiceForge supports the fitting of probability density data to the standard continuous probability distributions in section 6.

DiceForge also includes a ready-made random number generator `Random`, a global instance of the LFSR64 class. Users can run all the functions in section 2 on `Random` without delving into the object-oriented programming aspects.

For readability and convenience, DiceForge often uses the following predefined datatypes:

- `int32_t int64_t int128_t` Signed 32-bit, 64-bit, 128-bit integers
- `uint32_t uint64_t uint128_t` Unsigned 32-bit, 64-bit, 128-bit integers
- `real_t` A 64-bit floating point number (double)
- `int_t` A 64-bit signed integer
- `uint_t` A 64-bit unsigned integer

2 Functions of the RNGs

2.1 Bookkeeping functions

```
void rng.reseed(T seed)
```

Initializes the RNG with specified seed.

T is the data type supported by the derived class

If a is 0, the current system time is used as the seed. Otherwise, the integer a is used directly.

2.2 Functions to generate integers

```
T rng.next()
```

Generates and returns a uniformly selected random integer of type T (by default, unsigned long long int).

```
T rng.next_in_range(T min, T max)
```

Generates and returns a uniformly selected random integer of type T between min and max (both inclusive).

2.3 Functions to generate floats

```
real_t rng.next_unit()
```

Generates and returns a uniformly selected real number of type real_t between 0 (inclusive) and 1 (exclusive).

```
real_t rng.next_in_crange(real_t min, real_t max)
```

Generates and returns a uniformly selected random real number of type real_t between min (inclusive) and max (exclusive).

2.4 Functions acting on sequences

```
auto rng.choice(RandomAccessIterator first, RandomAccessIterator last)
```

Returns a randomly chosen element from the sequence defined by `first` and `last`.

```
void rng.shuffle(RandomAccessIterator first, RandomAccessIterator last)
```

Randomly shuffles the sequence defined by `first` and `last`, in place.

3 Functions of the Distributions

3.1 Continuous

```
real_t dist.expectation()
```

Returns the theoretically calculated expectation value of the distribution as a floating point number.

```
real_t dist.variance()
```

Returns the theoretically calculated variance of the distribution as a floating point number.

```
real_t dist.minValue()
```

Returns the theoretical minimum value of the distribution as a floating point number.

```
real_t dist.maxValue()
```

Returns the theoretical maximum value of the distribution as a floating point number.

```
real_t dist.pdf(real_t x)
```

Returns (as a floating point number) the probability density of the distribution at a floating point number x.

```
real_t dist.cdf(real_t x)
```

Returns (as a floating point number) the probability of generating a floating point number less than or equal to the floating point number x by the distribution.

3.2 Discrete

```
real_t dist.expectation()
```

Returns the theoretically calculated expectation value of the distribution as a floating point number.

```
real_t dist.variance()
```

Returns the theoretically calculated variance of the distribution as a floating point number.

```
int dist.minValue()
```

Returns the theoretical minimum value of the distribution as an integer.

```
int dist.maxValue()
```

Returns the theoretical maximum value of the distribution as an integer.

```
real_t dist.pmf(int x)
```

Returns (as a floating point number) the probability of generating the integer x by the distribution.

```
real_t dist.cdf(int x)
```

Returns (as a floating point number) the probability of generating an integer less than or equal to the integer x by the distribution.

4 Functions for Fitting Data

DiceForge also supports fitting a given set of points representing the probability density function of a distribution and fits it to one of the standard **continuous** distributions enlisted in section 6.

4.1 Fitting to a Cauchy distribution

```
DiceForge::Cauchy fitToCauchy(vector x, vector y, int max_iter,  
real_t epsilon)
```

Fits the given sample points ($x, y=\text{pdf}(x)$) to a Cauchy distribution using non-linear least squares regression.

Here **x** is the list of x coordinates as a `std::vector<real_t>` and **y** is the list of the corresponding y coordinates as a `std::vector<real_t>`. **max_iter** is the maximum iterations to attempt to fit the data (higher to try for better fits) and **epsilon** is the minimum acceptable error tolerance while attempting to fit the data (smaller to try for better fits).

The function returns a `DiceForge::Cauchy` distribution fit to the given sample points

4.2 Fitting to an Exponential distribution

```
DiceForge::Exponential fitToExponential(std::vector<real_t> x,  
std::vector<real_t> y, int max_iter, real_t epsilon)
```

Fits the given sample points ($x, y=\text{pdf}(x)$) to an Exponential distribution using a variant of Stochastic Gradient Descent (SGD) with a learning rate that decreases over time. The algorithm takes the natural logarithm of the pdf and optimizes using SGD with the mean square error of the cost function.

Here **x** is the list of x coordinates as a `std::vector<real_t>` and **y** is the list of the corresponding y coordinates as a `std::vector<real_t>`, **max_iter** is the maximum iterations to attempt to fit the data and **epsilon** is the minimum acceptable error tolerance while attempting to fit the data.

(defaults: **max_iter** = 10000, **epsilon** = 1e-6, optimum value in general cases, change according to requirement)

The function returns a `DiceForge::Exponential` distribution fit to the given sample points.

4.3 Fitting to an Gaussian distribution

```
DiceForge::Cauchy fitToGaussian(vector x, vector y, int max_iter,  
real_t epsilon)
```

Fits the given sample points ($x, y=\text{pdf}(x)$) to a Gaussian distribution using non-linear least squares regression.

Here **x** is the list of x coordinates as a `std::vector<real_t>` and **y** is the list of the

corresponding y coordinates as a `std::vector<real_t>`. `max_iter` is the maximum iterations to attempt to fit the data (higher to try for better fits) and `epsilon` is the minimum acceptable error tolerance while attempting to fit the data (smaller to try for better fits).

The function returns a `DiceForge::Gaussian` distribution fit to the given sample points

4.4 Fitting to a Maxwell distribution

```
DiceForge::Maxwell fitToMaxwell(vector x, vector y, int max_iter,  
real_t epsilon)
```

Fits the given sample points ($x, y=\text{pdf}(x)$) to a Maxwell distribution using non-linear least squares regression.

Here `x` is the list of x coordinates as a `std::vector<real_t>` and `y` is the list of the corresponding y coordinates as a `std::vector<real_t>`. `max_iter` is the maximum iterations to attempt to fit the data (higher to try for better fits) and `epsilon` is the minimum acceptable error tolerance while attempting to fit the data (smaller to try for better fits).

The function returns a `DiceForge::Maxwell` distribution fit to the given sample points.

4.5 Fitting to a Weibull distribution

```
DiceForge::Weibull fitToWeibull(vector x, vector y, int max_iter,  
real_t epsilon)
```

Fits the given sample points ($x, y=\text{pdf}(x)$) to a Weibull distribution using non-linear least squares regression.

It approximates the initial guess by calculating CDF and applying regression to find the parameters λ and k . The regression of CDF method gives very good guess for initial starting point for the Gauss-Newton method, as the CDF of Weibull is a linearizable function. After the guess is found, the non-linear least squares regression is used to approach and fine tune the final parameter values.

Here `x` is the list of x coordinates as a `std::vector<real_t>` and `y` is the list of the corresponding y coordinates as a `std::vector<real_t>`. `max_iter` is the maximum iterations to attempt to fit the data (higher to try for better fits) and `epsilon` is the minimum acceptable error tolerance while attempting to fit the data (smaller to try for better fits).

The function returns a `DiceForge::Weibull` distribution fit to the given sample points.

5 Generators supported

5.1 Linear Feedback Shift Register

A linear feedback shift register (LFSR) is an algorithm which generates psuedo-random numbers by performing certain linear operations on the seed. The DiceForge library implements a specific type of LFSR called XORshift LFSR, where the linear operation is the exclusive OR. To introduce non-linearity (and thus decrease predictability), the generated number is multiplied with a constant before outputting.

5.1.1 Algorithm

1. Define a function $f_{p,q,r,s}(x)$ over $x \in \mathbb{N}$, such that $f_{p,q,r,s}(x) = ((x \gg p) \oplus (x \gg q) \oplus (x \gg r) \oplus (x \gg s)) \& 1$. This function outputs a single bit, which will be pushed to the MSB.
2. Choose a seed $s (\neq 0)$ and set $x_0 = s$.
3. Choose a constant M and natural numbers p, q, r and s . p, q, r and s are chosen so as to get the maximum period of the LFSR.
4. Choose the required number of bits, N , in the output.
5. Initialise the output y_0 to zero.
6. Define a function $g_l(x, b)$ over $x \in \mathbb{N}$ and $b \in \{0, 1\}$, such that $g_l(x, b) = (x \gg 1) \& (b < (l - 1))$. Here l is the number of bits in the seed s .
7. $b = f_{p,q,r,s}(x_n)$
 $x_{n+1} = g_l(x_n, b)$
 $y_{n+1} = 2y_n + (x_{n+1} \& 1)$
8. The sequence of random numbers obtained from the states $\{x_n\}$ is $\{y_{64n} \times M\}$ for $n \in \mathbb{N}$.

DiceForge implements this algorithm with the following values:

- $l = 128$.
- $M = (2545F4914F6CDD1D)_{16}$.
- $\langle p, q, r, s \rangle = \langle 0, 1, 2, 7 \rangle$. This particular tuple gives a cycle length of $2^{128} - 1$, which is the highest possible for a 128-bit seed.
- $N = 64$ (for LFSR64) or 32 (for LFSR32).

5.2 Blum Blum Shub

The Blum-Blum-Shub follows a deterministic algorithm to output a sequence that appears to be random. That is, no poly-time algorithm will be able to distinguish between an output sequence of the BBS generator and a truly random sequence with probability significantly greater than $1/2[1]$. It is widely used for Cryptographic purposes and the security of BBS generator is as hard as factoring a large composite integer.

5.2.1 Algorithm

1. Select 2 sufficiently large prime numbers namely p and q , such that $p \equiv q \equiv 3 \pmod{4}$
2. Let, $n = p \times q$, also known as the *blum integer*
3. Choose a seed s , such that $1 \leq s \leq n - 1$ and $\gcd(s, n) = 1$
4. $x_0 = s^2$
5. Sequence is defined as: $x_i = x_{i-1}^2 \pmod{n}$

The algorithm operates on a seed value to produce an output, where each output serves as the seed for the subsequent iteration. The least significant bit of each output is extracted and concatenated through left shifting. This iterative process continues until a number within the range specified by the user is obtained.

Note:

$p = 60539$; $q = 50147$;

- Values of p and q chosen are safe primes which make the chain length larger.
- The chosen $n = p \times q$ is less than the square root of the largest number in unsigned 64bit long int; hence no overflow issues will occur even when seed is squared.
- In case, the algorithm reaches 1 or 0, a random number is taken as the seed so as to prevent the algorithm from getting stuck at 1 or 0.

5.3 Mersenne Twister

Implemented with the Mersenne Twister algorithm, this PRNG is a potent pseudo-random number generator, with a large period of $2^{19937} - 1$ (a Mersenne prime). It excels in distributing 623 values with 32-bit precision, utilizing 624 words for the former and 312 words for 64-bit numbers. The adjusted parameterization for 32-bit and 64-bit integer generation is noteworthy. Although there is another algorithm which generates more random, the basic version has been implemented here. Refer the appendix for the parameters used.

5.3.1 Algorithm

We know that MT generates random numbers from 0 to $2^w - 1$ where w is number of bits in each number. This is a linear recurrence PRNG which is based on the relation

$$\mathbf{x}_{k+n} = \mathbf{x}_{k+m} \oplus (\mathbf{x}_k^u | \mathbf{x}_{k+1}^l)A, \quad (k = 0, 1, \dots)$$

In MT, we generate random numbers using a state vector of N seeds, i.e. we give $\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{n-1}$ as seeds. m is an arbitrary integer between 1 and n . We choose m approximately to be the mid point of the array. A parameter r ($0 \leq r \leq w - 1$) is also included in the recurrence. \mathbf{x}_k^u on the right hand side means the upper $w - r$ bits of \mathbf{x}_k and similarly \mathbf{x}_{k+1}^l means the lower r bits of \mathbf{x}_{k+1} . Thus $(\mathbf{x}_k^u | \mathbf{x}_{k+1}^l)$ is concatenation of first $w - r$ bits of \mathbf{x}_k and last r bits of \mathbf{x}_{k+1} in that order. Then the matrix A is post-multiplied by this vector. Finally, XOR operation is performed with the vector \mathbf{x}_{k+m} to obtain the next number \mathbf{x}_{k+n} .

A is chosen so that multiplication with the vector becomes easier:

$$A = \begin{pmatrix} 0 & 1 & \dots & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ 0 & 0 & 0 & \ddots & \\ 0 & 0 & 0 & 0 & 1 \\ a_{w-1} & a_{w-2} & \dots & \dots & a_0 \end{pmatrix}$$

Now multiplication with A will be simplified to carrying the following operation:

$$\begin{cases} \mathbf{x} \gg 1 & \text{if } x_0 = 0 \\ (\mathbf{x} \gg 1) \oplus \mathbf{a} & \text{if } x_0 = 1 \end{cases}$$

where $\mathbf{a} = (a_{w-1}, a_{w-2}, \dots, a_0)$ and $\mathbf{x} = (x_{w-1}, x_{w-2}, \dots, x_0)$. AND operation with a suitable number will give first $w - r$ bits and r bits of \mathbf{x}_k and \mathbf{x}_{k+1} respectively

To make the distribution to k -distribution up to v -bit accuracy, we multiply the generated number with a suitable $w \times w$ invertible matrix T . The operation is the same as performing the following transformations.

$$\begin{aligned} \mathbf{y} &:= \mathbf{x} \oplus (\mathbf{x} \gg u) \\ \mathbf{y} &:= \mathbf{y} \oplus ((\mathbf{y} \ll s) \& \mathbf{b}) \\ \mathbf{y} &:= \mathbf{y} \oplus ((\mathbf{y} \ll t) \& \mathbf{c}) \\ \mathbf{z} &:= \mathbf{y} \oplus (\mathbf{x} \gg l) \end{aligned}$$

where l, s, u and t are suitable numbers, \mathbf{b} and \mathbf{c} are bit masks of length w , and z is the required random number.

5.4 XOR Shift

XOR Shift is a computationally inexpensive and memory compact algorithm for generating pseudo-random numbers. DiceForge presently follows the naive implementation of the XOR Shift algorithm followed by a non-linear multiplicative transform for improved quality.

5.4.1 Algorithm

The algorithm can be stated as follows:

1. We will first define a function $S_{p,q,r}(x)$ where $x \in \mathbb{N}$.
Let $l_1 = x \oplus (x \ll p)$, and $l_2 = l_1 \oplus (l_1 \gg q)$. Then, $S_{p,q,r}(x) = l_2 \oplus (l_2 \ll r)$
Here $(A \gg k)$ denotes the number obtained by shifting the binary representation of A by k bits to the right and setting all empty bits in the left end to be zero and $(A \ll k)$ denotes the number obtained by shifting the binary representation of A by k bits to the left and all setting all empty bits in the right end to be zero
2. Choose the tuple $\langle p, q, r \rangle$ and the constant M to be used for generating the pseudo-random state sequence $\{x_n\}$
3. Take an initial seed $s (\neq 0)$ and set $x_0 = s$.
4. $x_{n+1} = S_{p,q,r}(x_n)$ defines the relation between the states for $n \in \mathbb{N}$
and the sequence of random numbers obtained from these states is $\{M \times x_n\}$

The implementation of the algorithm in DiceForge is with,

- the following choice of $\langle p, q, r \rangle$ from the original paper by Marsaglia to ensure high periods of $2^{32} - 1$ and $2^{64} - 1$ respectively: $\langle 13, 17, 5 \rangle$ for the 32-bit implementation and $\langle 13, 7, 17 \rangle$ for the 64-bit implementation
- $M = (2545F4914F6CDD1D)_{16}$

5.5 Naor-Reingold Pseudo-random Function

Noar Reingold is an efficient pseudorandom function used for various cryptographic purposes like symmetric encryption, authentication and digital signatures. The security of the function lies in the difficulty in predicting the function value $f_a(x)$ when the attacker knows some terms of the function $f_a(x)$.

5.5.1 Algorithm

1. Let p and q be two prime numbers such that $l|p-1$.
2. Select an element $g \in \mathbb{F}_p^*$ of multiplicative order l , simply $g^l \equiv 1 \pmod{p}$.
3. For a particular n , define a $(n+1)$ dimensional vector $a = (a_0, a_1, \dots, a_n) \in \mathbb{F}_l^{n+1}$.
4. Then the pseudorandom number generated is

$$f_a(x) = g^{a_0.a_1^{x_1}.a_2^{x_2}....a_n^{x_n}} \pmod{p} \quad (1)$$

where $x = (x_1, x_2, \dots, x_n)$ is the bit representation of initial seed value x .

While implementing the algorithm in DiceForge the following values of the variables have been used:

- $p = 4279969613$ and $l = 9999929$
- $g = 9999918$
- $n = 32$

6 Continuous distributions supported

6.1 Cauchy

6.1.1 Description

The Cauchy Distribution, also known as the Lorentzian or Cauchy-Lorentzian distribution is a normal distribution that often finds applications in physical and mathematical modelling. A Cauchy distribution is governed by two parameters x_0 and γ . An intuitive way to think of the distribution is to consider the following problem.

In the 2D plane, shoot rays originating from (x_0, γ) making an angle θ with some fixed reference axis. If θ is uniformly distributed, then if a random variable X takes the values of the x-intercept made by the rays, then X follows a Cauchy distribution. The probability density of this distribution is given by:

$$f(x) = \frac{1}{\pi\gamma \left[1 + \left(\frac{x-x_0}{\gamma}\right)^2\right]} \quad (2)$$

6.1.2 Parameters

- $x_0 > 0$: Shift parameter (determines the centre of the distribution)
- $\gamma > 0$: Scale parameter (determines the spread of the distribution)

6.1.3 Properties

The Cauchy distribution is a probability distribution describing a continuous random variable whose moments of integer order greater than zero do not exist. So, the expectation and variance of a Cauchy distribution are undefined.

6.1.4 Usage

This distribution is modelled in DiceForge by the class `DiceForge::Cauchy` which is derived from the base `DiceForge::Continuous`. A distribution can be easily created by instantiating the class and feeding in the parameters of the distribution to the constructor `DiceForge::Cauchy(real_t x0, real_t gamma)` and using the various methods of the base class `DiceForge::Continuous` as specified in section 3.1.

6.2 Exponential

6.2.1 Description

The exponential distribution is a continuous probability distribution commonly used to model the time until an event occurs. It is characterized by a single parameter k , representing the rate parameter, where $k > 0$, and a change of origin x_0 .

The probability density function (PDF) for a random variable X following an exponential distribution is given by:

$$P(X = x) = \begin{cases} ke^{-k(x-x_0)}, & \text{if } x \geq x_0 \\ 0, & \text{otherwise} \end{cases} \quad (3)$$

6.2.2 Parameters

- $k > 0$: Rate parameter (determines the rate of the distribution).
- x_0 : Change of origin (default $x_0 = 0$).

6.2.3 Properties

- The mean of the Exponential distribution is given by $E[X] = \frac{1}{k} + x_0$.
- The variance of the Exponential distribution is given by $\text{Var}[X] = \frac{1}{k^2}$.

6.2.4 Usage

This distribution is modelled in DiceForge by the class `DiceForge::Exponential`, which is derived from the base class `DiceForge::Continuous`. To instantiate this, you need to provide the rate parameter k and the change of origin x_0 to the constructor using

Class Constructor:

```
DiceForge::Exponential(real_t k, real_t x0)
```

The constructor initialises the distribution with parameters detailed in the previous sections. The various methods of the base class `DiceForge::Continuous` can be used as specified in section 3.1.

6.3 Gaussian (Normal)

6.3.1 Description

The Gaussian or Normal distribution is a continuous probability distribution symmetric about its mean. It is characterised by two parameters μ and σ , where μ is the mean of the distribution and σ is the standard deviation.

The probability of the random variable X for a normal distribution is given by:

$$P(X = x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (4)$$

6.3.2 Parameters

- μ : Shift parameter (determines the centre of the distribution)
- $\sigma > 0$: Scale parameter (determines the spread of the distribution)

6.3.3 Properties

- The mean of the Gaussian distribution is given by $E[X] = \mu$.
- The variance of the Maxwell distribution is given by $\text{Var}[X] = \sigma$.
- It is a unimodal distribution with the peak at its mean $x = \sigma$.
- About 68% of the random variable falls within one standard deviation, σ , from the mean, about 95% within two standard deviations, and about 99.7% within three standard deviations.

6.3.4 Usage

This distribution is modelled in DiceForge by the class `DiceForge::Gaussian` which is derived from the base `DiceForge::Continuous`. A distribution can be easily created by instantiating the class and feeding in the parameters of the distribution to the constructor `DiceForge::Gaussian(real_t mu, real_t sigma)` and using the various methods of the base class `DiceForge::Continuous` as specified in section 3.1.

6.4 Maxwell-Boltzmann

6.4.1 Description

The Maxwell-Boltzmann distribution is a probability distribution describing the speeds of the particles in idealized gases. It's characterised by a single parameter $a = \sqrt{\frac{kT}{m}}$, where k is the Boltzmann constant, T is the temperature, m is the molecular mass of the gas particle.

The probability density function of this distribution is given by:

$$f(x) = \sqrt{\frac{2}{\pi}} \frac{x^2}{a^3} \exp\left(\frac{-x^2}{2a^2}\right) \quad (5)$$

6.4.2 Parameters

- $a > 0$: Scale parameter (determines the spread of the distribution)

6.4.3 Properties

- The mean of the Maxwell distribution is given by $E[X] = 2a\sqrt{\frac{2}{\pi}}$.
- The variance of the Maxwell distribution is given by $Var[X] = \frac{a^2(3\pi-8)}{\pi}$.

6.4.4 Usage

This distribution is modelled in DiceForge by the class `DiceForge::Maxwell` which is derived from the base `DiceForge::Continuous`. A distribution can be easily created by instantiating the class and feeding in the parameters of the distribution to the constructor `DiceForge::Maxwell(real_t a)` and using the various methods of the base class `DiceForge::Continuous` as specified in section 3.1.

6.5 Weibull

6.5.1 Description

The Weibull distribution is a continuous probability distribution that is often used in reliability engineering, life data analysis, and survival analysis. It offers flexibility in modeling various types of data with different shapes and scales.

It is characterized by its shape parameter k and scale parameter λ .

The probability density function (pdf) and cumulative distribution function (cdf) of the Weibull distribution are given by:

$$f(x; k, \lambda) = \begin{cases} \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} e^{-\left(\frac{x}{\lambda}\right)^k}, & x > 0 \\ 0, & \text{otherwise} \end{cases} \quad (6)$$

$$F(x; k, \lambda) = 1 - e^{-\left(\frac{x}{\lambda}\right)^k} \quad (7)$$

6.5.2 Parameters

- $k > 0$: Shape parameter (determines the shape of the distribution).
- $\lambda > 0$: Scale parameter (determines the scale or location of the distribution).

6.5.3 Properties

- The mean of the Weibull distribution is given by $E[X] = \lambda \Gamma\left(1 + \frac{1}{k}\right)$, where $\Gamma(\cdot)$ is the gamma function.
- The variance of the Weibull distribution is given by $\text{Var}[X] = \lambda^2 \left[\Gamma\left(1 + \frac{2}{k}\right) - \left(\Gamma\left(1 + \frac{1}{k}\right)\right)^2 \right]$.

6.5.4 Usage

Class Constructor:

```
DiceForge::Weibull(real_t lambda, real_t k)
```

The constructor initialises the distribution with parameters detailed in the previous sections.

6.6 Custom Distribution

6.6.1 Description

The Custom Distribution allows the user to define a custom probability density function(PDF) and sample data points according to the provided PDF. The class empowers you to experiment with novel probability distributions for research purposes. You can define theoretical PDF's and use the sample method to generate data for analysis and hypothesis testing.

6.6.2 Parameters

- PDF that defines the distribution: The Custom Distribution class relies on the user to provide a valid probability density function.

Check-points for a valid PDF:

1. Non-negativity: The PDF ($f(x)$) must always return non-negative values for all possible values of x within the defined range.

$$P(X = x) = f(x) \geq 0 \quad (8)$$

for all relevant x values.

2. Integrability over the Domain: The total area under the PDF curve, over the specified domain, must integrate to 1.

$$\int_a^b f(x) dx = 1 \quad (9)$$

- Domain of the PDF (upper and lower limits of integration)

6.6.3 Properties

All properties of the distribution are calculated numerically. Existence of these properties does not imply integrability of the function. The domain over which the function has to be evaluated is broken down into smaller intervals dictated by a step size h . The step size is currently set to 0.001, and can be altered depending on the required precision of sampling. You can achieve a more precise result by lowering the value, but this will come at the cost of increased computation time. Methods by which these properties are calculated are discussed below

- Expectation: $x_0 = \sum x * f(x) * h$
- Variance: $\sigma^2 = \sum (x - x_0)^2 * f(x) * h$
- Cumulative density function(CDF): The CDF of the function is calculated at each interval using Simpson's method of integration. For a specific value of x , linear interpolation is used to estimate the CDF based on the two encompassing pre-computed CDF values from the nearest intervals.

6.6.4 Usage

This class can be initialised by the constructor:

```
DiceForge::CustomDistribution(real_t lower, real_t upper, PDF_Function pdf) . where  
PDF_Function is defined as follows PDF_Function = std::function<real_t(real_t)>
```

7 Discrete distributions supported

7.1 Bernoulli

7.1.1 Description

A Bernoulli distribution is used to model the outcome of a random experiment that has only two possible outcomes: success (usually denoted as 1) or failure (usually denoted as 0).

$$P(X = x) = \begin{cases} p & \text{if } x = 1 \\ 1 - p & \text{if } x = 0 \end{cases} \quad (10)$$

7.1.2 Parameters

- p , the probability of a success.

7.1.3 Usage

```
DiceForge::Bernoulli (real_t p)
```

To instantiate this distribution, you need to provide the parameter p to the constructor.

7.2 Binomial

7.2.1 Description

A binomial distribution is a probability distribution that describes the number of successes in a fixed number of independent Bernoulli trials, where each trial has only two possible outcomes: success or failure. The probability mass function (PMF) of a binomial distribution, denoted as $P(X = k)$, gives the probability of observing exactly k successes out of n trials. It is given by the formula:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{(n-k)}$$

The mean is given by $\mu = np$, and the variance is given by $\sigma^2 = np(1 - p)$

7.2.2 Parameters

- n , the number of trials.
- p , the probability of each trial being a success.

7.2.3 Usage

```
DiceForge::Binomial (uint_t n, real_t p)
```

The class constructor takes in two parameters, `n`, the number of trials (an unsigned integer), and `p`, the probability of success (a float between 0 and 1, inclusive). It creates and stores the probability mass function on initialization. Hence, it requires $O(n)$ preprocessing time, and $O(1)$ for each subsequent query.

7.3 Geometric

7.3.1 Definition

The geometric distribution is either one of two discrete probability distributions:

- The probability distribution of the number X of Bernoulli trials needed to get one success, supported on the set $\{1, 2, 3, \dots\}$;
- The probability distribution of the number $Y = X - 1$ of failures before the first success, supported on the set $\{0, 1, 2, 3, \dots\}$;

Which of these is called the geometric distribution is a matter of convention and convenience. (Here we've considered the first definition)

7.3.2 Parameters and PMF

The geometric distribution gives the probability that the first occurrence of success requires k independent trials, each with success probability p . If the probability of success on each trial is p , then the probability that the k -th trial is the first success is-

$$Pr(X = k) = p(1 - p)^{k-1}; (0 < p \leq 1)$$

7.3.3 Properties

The expected value for the number of independent trials to get the first success, and the variance of a geometrically distributed random variable X is:

$$E(X) = 1/p$$

$$var(X) = (1 - p)/p^2$$

7.3.4 Usage

```
DiceForge::Geometric (real_t p)
```

In DiceForge, the geometric distribution is implemented by the class `DiceForge::Geometric`, which is derived from the base class `DiceForge::Discrete`. To instantiate this distribution, you need to provide the parameter p to the constructor `DiceForge::Geometric(real_t p)`

7.4 Gibbs

7.4.1 Definition and applications

The Gibbs distribution is a discrete random variable distribution, primarily used in statistical mechanics as the Boltzmann distribution. In that context, it gives the probability that a system will be in a certain energy state at a given temperature. The probability of the random variable X having the value x (or the system X being in state x) is given by:

$$P(X = x) = \frac{1}{Z(\beta)} e^{-\beta E(x)} \quad (11)$$

where E is a function of x (in physics, it is the energy of state x), β is an arbitrary parameter (in physics it is the inverse of the absolute temperature), and $Z(\beta)$ is the normalising constant.

7.4.2 Parameters

- $E(x)$, a function on x .
- β , an arbitrary parameter.

7.4.3 Usage

```
DiceForge::Gibbs(RandomAccessIterator1 sequence_first, RandomAccessIterator1  
sequence_last, RandomAccessIterator2 function_first,  
RandomAccessIterator2 function_last, real_t beta)
```

The class constructor `Gibbs()` takes input `sequence_first` and `sequence_last` (a sequence of numerical x values for the distribution), `function_first` and `function_last` (a sequence of the corresponding numerical $E(x)$ values), and `beta` (the parameter). Creates and stores arrays corresponding to the PMF and CDF of the distribution. Also stores the sequence of x values and its length internally.

```
int_t DiceForge::Gibbs::next(real_t r)
```

The `next()` function takes as input a random floating-point number between 0 and 1, and outputs a corresponding integer chosen according to the Gibbs distribution.

7.5 Hypergeometric

7.5.1 Description

The hypergeometric distribution is a probability distribution that describes the number of successes (or items of a particular type) in a fixed-size sample drawn without replacement from a finite population containing a specific number of successes and failures. It is used when the outcomes are not independent, as the sampling is done without replacement.

7.5.2 Parameters

The parameters involved here are N, K, n where N is the size of the population, K members are of a particular type (which defines success) and n is the sample size. The distribution defines a random variable X which is equal to number of successes in drawing a sample of size n . Note that $n \geq 0, K \geq 0, n \leq N$ and $K \leq N$. so it throws an error when constraints are not satisfied.

7.5.3 Properties

The probability mass function is given by

$$P(X = x) = \frac{\binom{K}{x} \binom{N-K}{n-x}}{\binom{N}{n}}$$

The mean of the distribution is

$$\text{Mean} = n \frac{K}{N}$$

The variance of the distribution is

$$\text{Variance} = n \frac{K}{N} \frac{N-K}{N} \frac{N-K}{N-1}$$

7.5.4 Usage

```
DiceForge::Hypergeometric(int N, int K, int n)
```

The class constructor initialises a vector to store the probability mass function for each x recursively from $\max(n + K - N, 0)$ to $\min(n, K)$. using which CDF can be calculated and stored internally.

7.6 Negative Hypergeometric

7.6.1 Description

The negative hypergeometric distribution defines the behaviour of the following random variable X , governed by three parameters N , k and r .

Consider a population of size N whose members can be classified exhaustively into two mutually exclusive categories, A and \bar{A} where K members belong to category A and the remaining $N - K$ members belong to category \bar{A} . Now members of the population are randomly chosen (without replacement) and are identified as belonging to A or \bar{A} . Once exactly r members from category \bar{A} are selected this procedure is stopped. Suppose now a total of $r + k$ members from the population have been chosen, then we set $X = k$.

7.6.2 Properties

The probability of the random variable X taking the value k in this distribution is given by the PMF:

$$P(X = k) = \frac{\binom{k+r-1}{k} \binom{N-k-r}{K-k}}{\binom{N}{K}} \quad (12)$$

The expectation value and variance of the random variable X are given by:

$$E[X] = \frac{rK}{N - K + 1} \quad (13)$$

$$Var[X] = \frac{(N+1)rK}{(N-K+1)(N-K+2)} \left(1 - \frac{r}{N-K+1} \right) \quad (14)$$

7.6.3 Parameters

- N , size of the population
- K , number of success elements in the population
- r , number of failure elements to be picked from the population for the procedure to stop

7.6.4 Usage

This distribution is modelled in DiceForge by the class `DiceForge::NegHypergeometric` which is derived from the base `DiceForge::Discrete`. A distribution can be easily created by instantiating the class and feeding in the parameters of the distribution to the constructor `DiceForge::NegHypergeometric(uint N, uint K, uint r)` and using the various methods of the base class `DiceForge::Discrete` as specified in section 3.2

7.7 Poission

7.7.1 Definition

The Poission distribution is a discrete random variable distribution which represents the probability of a given number of events happening in a fixed interval of time given that these events occur with a fixed mean rate and the time period between successive events is independent of the previous event. It is characterized by the aforementioned mean rate λ . Given λ , the probability of k events occurring is given by:

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

The CDF of the distribution does not have a closed form with respect to normal operations but may be represented in terms of the incomplete gamma function:

$$P(X \leq k) = \frac{\Gamma(k+1, \lambda)}{\Gamma(k+1)}$$

7.7.2 Implementation

In DiceForge, the numbers are generated by employment of the rejection method for discrete distributions. A curve with an invertible PDF expression is first chosen such that it is greater than the Poisson curve at all points. Random numbers are then generated such that they are uniform in the area underneath the curve. If a generated random number also lies under the Poisson curve, it is accepted. Otherwise, it is rejected.

7.7.3 Parameters

- λ , characterises the distribution as mentioned in [7.7.1](#)

7.7.4 Usage

The distribution is implemented using a class `DiceForge::Poission` and may be instantiated as an object (say `poisson_object`) by the constructor:

```
DiceForge::Poission(real_t lambda)
```

While sampling random numbers using the instance of the class, you have to call:

`poisson_object.next(DiceForge::Generator& rng)` where `rng` is an object of any RNG class currently supported by the DiceForge library

8 2D Random Variables

This section includes the documentation for usage of functions for integrating one dimensional and two dimensional random variable probability distribution function over a specified region to obtain the probability of the random variable taking value in that region. Note that at least C++ version 20 is required to use these functions.

8.1 Integrate (1D) Function Usage

```
DiceForge::integrate<FuncType, BoundType>(FuncType f, std::tuple<BoundType,  
BoundType> bounds)
```

or put simply `integrate(function f(x), pair bounds)`

To use the `integrate` function, follow the template:

```
1 // Define your function to be integrated
2 real_t my_function(real_t x) {
3     // Define your function here
4     return x * x;
5 }
6 // Example usage of the integrals function
7 auto result = DiceForge::integrate(my_function, std::make_tuple(0.0, 1.0));
```

8.2 Integrate (2D) Function Usage

```
DiceForge::integrate(FuncType f, std::tuple<Lower_0, Upper_0> bound_0,  
std::tuple<Lower_1, Upper_1> bound_1, std::integer_sequence<int, First, Second>)
```

or put simply `integrate(function f(x, y), pair bound0, pair bound1, integration_sequence)`

To use the `integrate` function, follow the template:

```
1 // Define your function to be integrated
2 real_t my_function(real_t x, real_t y) {
3     // Define your function here
4     return x + y * y;
5 }
6 // Example usage of the integrals function
7 auto result = DiceForge::integrate(my_function, std::make_tuple(0.0, 1.0),  
std::make_tuple(-1.0, 1.0), DiceForge::dx_dy);
```

Here the function type and bound types `lower_0` and `upper_0` are detected automatically by compiler by template auto deduction and hence there's no need to actually specify them. Note that the bounds_1 can be functions of y (if dx dy is the order). This can be done by using function pointers or using lambdas. Also, for sequence as a parameter in the function, use

`DiceForge::dx_dy` – for integration order being x then y

`DiceForge::dy_dx` – for integration order being y then x

9 Performance and Statistics

9.1 Uniformity of the PRNGs

An analysis for the uniformity of the PRNGs using mean and variance tests over 10^8 numbers yields the following results:

Generator	Mean	Variance
BBS32	0.501565	0.0830039
BBS64	0.501565	0.0830038
XORShift32	0.499978	0.08334
XORShift64	0.499931	0.0833327
MT32	0.499956	0.0833425
MT64	0.500034	0.0833241
LFSR32	0.500006	0.0833379
LFSR64	0.499971	0.0833363
NaorReingold	0.49824	0.0827696

Table 1: Uniformity of the Generators

For reference, an ideal uniform distribution should have a mean of 0.5 and variance of $\frac{1}{12} = 0.08\bar{3}$.

9.2 Time Performance

Generator	Time Taken
BBS32	15210.2ms
BBS64	24677.3ms
XORShift32	484.018ms
XORShift64	441.947ms
MT32	1742.37ms
MT64	1751.81ms
LFSR32	10000.9ms
LFSR64	18665.6ms
NaorReingold	195678ms

Table 2: Time taken to generate 10^8 random numbers by DiceForge

Generator	Time Taken
C++'s Mersenne Twister	4032.7ms
C rand() function	1388.03ms
python's random	175175.89ms ($\approx 3\text{min}$)
numpy's randint	165600.48ms ($\approx 3\text{min}$)

Table 3: Time taken to generate 10^8 random numbers by other libraries

9.3 Dieharder Tests Results

9.3.1 Mersenne Twister

32 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.96960356	PASSED
diehard_operm5	1000000	100	0.02305118	PASSED
diehard_rank_32x32	40000	100	0.80573118	PASSED
diehard_rank_6x8	100000	100	0.29893285	PASSED
diehard_bitstream	2097152	100	0.97759639	PASSED
diehard_opso	2097152	100	0.99911115	WEAK
diehard_oqso	2097152	100	0.99025734	PASSED
diehard_dna	2097152	100	0.15367448	PASSED
diehard_count_1s_stream	256000	100	0.93442989	PASSED
diehard_count_1s_byte	256000	100	0.99877	WEAK
diehard_parking_lot	12000	100	0.43795973	PASSED
diehard_2dsphere	8000	100	0.02227018	PASSED
diehard_3dsphere	4000	100	0.06358327	PASSED
diehard_squeeze	100000	100	0.99979944	WEAK
diehard_sums	100	100	0.60721104	PASSED
diehard_runs	100000	100	0.72693794	PASSED
diehard_runs	100000	100	0.17953541	PASSED
diehard_craps	200000	100	0.97986466	PASSED
diehard_craps	200000	100	0.94227466	PASSED
marsaglia_tsang_gcd	10000000	100	0.44019392	PASSED
marsaglia_tsang_gcd	10000000	100	0.63446613	PASSED
sts_monobit	100000	100	0.68004027	PASSED
sts_runs	100000	100	0.74546621	PASSED
sts_serial	100000	100	0.68004027	PASSED
sts_serial	100000	100	0.23264792	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.69887571	PASSED
rgb_lagged_sum	1000000	100	0.96017704	PASSED
rgb_kstest_test	10000	1000	0.32823119	PASSED
dab_bytedistrib	51200000	1	0.51605793	PASSED
dab_dct	50000	1	0.0683238	PASSED
dab_filltree	15000000	1	0.21923772	PASSED
dab_filltree	15000000	1	0.99522766	WEAK
dab_filltree2	5000000	1	0.37167458	PASSED
dab_filltree2	5000000	1	0.67928699	PASSED
dab_monobit2	65000000	1	0.94981924	PASSED

64 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.97296453	PASSED
diehard_operm5	1000000	100	0.87258133	PASSED
diehard_rank_32x32	40000	100	0.59176364	PASSED
diehard_rank_6x8	100000	100	0.95929486	PASSED
diehard_bitstream	2097152	100	0.49715412	PASSED
diehard_opso	2097152	100	0.99180199	PASSED
diehard_oqso	2097152	100	0.59827403	PASSED
diehard_dna	2097152	100	0.7642168	PASSED
diehard_count_1s_stream	256000	100	0.35158926	PASSED
diehard_count_1s_byte	256000	100	0.98892266	PASSED
diehard_parking_lot	12000	100	0.06301159	PASSED
diehard_2dsphere	8000	100	0.75428486	PASSED
diehard_3dsphere	4000	100	0.96001921	PASSED
diehard_squeeze	100000	100	0.74761552	PASSED
diehard_sums	100	100	0.00071988	WEAK
diehard_runs	100000	100	0.46378537	PASSED
diehard_runs	100000	100	0.67339262	PASSED
diehard_craps	200000	100	0.36710702	PASSED
diehard_craps	200000	100	0.55937118	PASSED
marsaglia_tsang_gcd	10000000	100	0.94291467	PASSED
marsaglia_tsang_gcd	10000000	100	0.93107236	PASSED
sts_monobit	100000	100	0.38393537	PASSED
sts_runs	100000	100	0.23249519	PASSED
sts_serial	100000	100	0.38393537	PASSED
sts_serial	100000	100	0.6152597	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.49066069	PASSED
rgb_lagged_sum	1000000	100	0.59590977	PASSED
rgb_kstest_test	10000	1000	0.19617897	PASSED
dab_bytedistrib	51200000	1	0.75150579	PASSED
dab_dct	50000	1	0.84885855	PASSED
dab_filltree	15000000	1	0.66993604	PASSED
dab_filltree	15000000	1	0.70370342	PASSED
dab_filltree2	5000000	1	0.34129307	PASSED
dab_filltree2	5000000	1	0.26522698	PASSED
dab_monobit2	65000000	1	0.87840968	PASSED

9.3.2 XORShift

32 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.39782176	PASSED
diehard_operm5	1000000	100	0.25255851	PASSED
diehard_rank_32x32	40000	100	0.975651	PASSED
diehard_rank_6x8	100000	100	0.54768254	PASSED
diehard_bitstream	2097152	100	0.87561862	PASSED
diehard_opso	2097152	100	0.85153201	PASSED
diehard_oqso	2097152	100	0.89389646	PASSED
diehard_dna	2097152	100	0.24625874	PASSED
diehard_count_1s_stream	256000	100	0.96670825	PASSED
diehard_count_1s_byte	256000	100	0.68903363	PASSED
diehard_parking_lot	12000	100	0.39609831	PASSED
diehard_2dsphere	8000	100	0.75981652	PASSED
diehard_3dsphere	4000	100	0.26384334	PASSED
diehard_squeeze	100000	100	0.47290812	PASSED
diehard_sums	100	100	0.07056891	PASSED
diehard_runs	100000	100	0.89971161	PASSED
diehard_runs	100000	100	0.86649512	PASSED
diehard_craps	200000	100	0.57206057	PASSED
diehard_craps	200000	100	0.26981857	PASSED
marsaglia_tsang_gcd	10000000	100	0.15249939	PASSED
marsaglia_tsang_gcd	10000000	100	0.29214027	PASSED
sts_monobit	100000	100	0.77198832	PASSED
sts_runs	100000	100	0.81560493	PASSED
sts_serial	100000	100	0.77198832	PASSED
sts_serial	100000	100	0.43860458	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.50166693	PASSED
rgb_lagged_sum	1000000	100	0.78860934	PASSED
rgb_kstest_test	10000	1000	0.99250122	PASSED
dab_bytedistrib	51200000	1	0.40882691	PASSED
dab_dct	50000	1	0.19608202	PASSED
dab_filltree	15000000	1	0.76253042	PASSED
dab_filltree	15000000	1	0.58088234	PASSED
dab_filltree2	5000000	1	0.28117471	PASSED
dab_filltree2	5000000	1	0.64785586	PASSED
dab_monobit2	65000000	1	0.00667117	PASSED

64 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.8692515	PASSED
diehard_operm5	1000000	100	0.94222658	PASSED
diehard_rank_32x32	40000	100	0.32084984	PASSED
diehard_rank_6x8	100000	100	0.80736731	PASSED
diehard_bitstream	2097152	100	0.38808783	PASSED
diehard_opso	2097152	100	0.96518137	PASSED
diehard_oqso	2097152	100	0.45833456	PASSED
diehard_dna	2097152	100	0.62350735	PASSED
diehard_count_1s_stream	256000	100	0.50347078	PASSED
diehard_count_1s_byte	256000	100	0.98623369	PASSED
diehard_parking_lot	12000	100	0.81126942	PASSED
diehard_2dsphere	8000	100	0.27969302	PASSED
diehard_3dsphere	4000	100	0.12663541	PASSED
diehard_squeeze	100000	100	0.66649951	PASSED
diehard_sums	100	100	0.00099715	WEAK
diehard_runs	100000	100	0.97664751	PASSED
diehard_runs	100000	100	0.7936488	PASSED
diehard_craps	200000	100	0.4038087	PASSED
diehard_craps	200000	100	0.61756653	PASSED
marsaglia_tsang_gcd	10000000	100	0.59289793	PASSED
marsaglia_tsang_gcd	10000000	100	0.05434497	PASSED
sts_monobit	100000	100	0.02489182	PASSED
sts_runs	100000	100	0.75139208	PASSED
sts_serial	100000	100	0.02489182	PASSED
sts_serial	100000	100	0.36120012	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.94500538	PASSED
rgb_lagged_sum	1000000	100	0.84803641	PASSED
rgb_kstest_test	10000	1000	0.63595477	PASSED
dab_bytedistrib	51200000	1	0.59711739	PASSED
dab_dct	50000	1	0.52626023	PASSED
dab_filltree	15000000	1	0.06664793	PASSED
dab_filltree	15000000	1	0.70983336	PASSED
dab_filltree2	5000000	1	0.98500863	PASSED
dab_filltree2	5000000	1	0.86956032	PASSED
dab_monobit2	65000000	1	0.10268683	PASSED

9.3.3 LFSR

32 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.38964691	PASSED
diehard_operm5	1000000	100	0.88140385	PASSED
diehard_rank_32x32	40000	100	0.69953861	PASSED
diehard_rank_6x8	100000	100	0.13280074	PASSED
diehard_bitstream	2097152	100	0.79718314	PASSED
diehard_opso	2097152	100	0.45033847	PASSED
diehard_oqso	2097152	100	0.34718577	PASSED
diehard_dna	2097152	100	0.19975594	PASSED
diehard_count_1s_stream	256000	100	0.60364122	PASSED
diehard_count_1s_byte	256000	100	0.99785125	WEAK
diehard_parking_lot	12000	100	0.70447638	PASSED
diehard_2dsphere	8000	100	0.51937166	PASSED
diehard_3dsphere	4000	100	0.51386416	PASSED
diehard_squeeze	100000	100	0.988034	PASSED
diehard_sums	100	100	0.9382803	PASSED
diehard_runs	100000	100	0.93791732	PASSED
diehard_runs	100000	100	0.82854192	PASSED
diehard_craps	200000	100	0.1537168	PASSED
diehard_craps	200000	100	0.93831803	PASSED
marsaglia_tsang_gcd	10000000	100	0.00476857	WEAK
marsaglia_tsang_gcd	10000000	100	0.94897284	PASSED
sts_monobit	100000	100	0.32007096	PASSED
sts_runs	100000	100	0.86717279	PASSED
sts_serial	100000	100	0.32007096	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.84196771	PASSED
rgb_lagged_sum	1000000	100	0.43701637	PASSED
rgb_kstest_test	10000	1000	0.81595686	PASSED
dab_bytedistrib	51200000	1	0.79232425	PASSED
dab_dct	50000	1	0.96652798	PASSED
dab_filltree	15000000	1	0.95729372	PASSED
dab_filltree	15000000	1	0.80050479	PASSED
dab_filltree2	5000000	1	0.68050409	PASSED
dab_filltree2	5000000	1	0.18776247	PASSED
dab_monobit2	65000000	1	0.08741817	PASSED

64 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.30889575	PASSED
diehard_operm5	1000000	100	0.47844582	PASSED
diehard_rank_32x32	40000	100	0.44396697	PASSED
diehard_rank_6x8	100000	100	0.66501057	PASSED
diehard_bitstream	2097152	100	0.70929518	PASSED
diehard_opso	2097152	100	0.98493625	PASSED
diehard_oqso	2097152	100	0.98517487	PASSED
diehard_dna	2097152	100	0.25160202	PASSED
diehard_count_1s_stream	256000	100	0.32623841	PASSED
diehard_count_1s_byte	256000	100	0.83509131	PASSED
diehard_parking_lot	12000	100	0.25249632	PASSED
diehard_2dsphere	8000	100	0.39042658	PASSED
diehard_3dsphere	4000	100	0.45950229	PASSED
diehard_squeeze	100000	100	0.87498307	PASSED
diehard_sums	100	100	0.19126483	PASSED
diehard_runs	100000	100	0.6956406	PASSED
diehard_runs	100000	100	0.60101553	PASSED
diehard_craps	200000	100	0.98553756	PASSED
diehard_craps	200000	100	0.05130232	PASSED
marsaglia_tsang_gcd	10000000	100	0.66640737	PASSED
marsaglia_tsang_gcd	10000000	100	0.97040518	PASSED
sts_monobit	100000	100	0.85541678	PASSED
sts_runs	100000	100	0.74743459	PASSED
sts_serial	100000	100	0.85541678	PASSED
sts_serial	100000	100	0.72207648	PASSED
rgb_minimum_distance	10000	1000	0	FAILED
rgb_permutations	100000	100	0.03499421	PASSED
rgb_lagged_sum	1000000	100	0.97267121	PASSED
rgb_kstest_test	10000	1000	0.38824957	PASSED
dab_bytedistrib	51200000	1	0.45673275	PASSED
dab_dct	50000	1	0.93034844	PASSED
dab_filltree	15000000	1	0.17327864	PASSED
dab_filltree	15000000	1	0.23889173	PASSED
dab_filltree2	5000000	1	0.95968732	PASSED
dab_filltree2	5000000	1	0.05485473	PASSED
dab_monobit2	65000000	1	0.01438031	PASSED

9.3.4 Naor Reingold (32 bit only)

Test Name	tsamples	psamples	p value	Result
diehard_birthdays	100	100	0.58768330	PASSED
diehard_operm5	1000000	100	0.96535968	PASSED
diehard_rank_32x32	40000	100	0.22702655	PASSED
diehard_rank_6x8	100000	100	0.60553642	PASSED
diehard_bitstream	2097152	100	0.00000017	FAILED
diehard_opso	2097152	100	0.00000000	FAILED
diehard_oqso	2097152	100	0.00000000	FAILED
diehard_dna	2097152	100	0.03809661	PASSED
diehard_count_1s_stream	256000	100	0.47513658	PASSED
diehard_count_1s_byte	256000	100	0.06509727	PASSED
diehard_parking_lot	12000	100	0.00000003	FAILED
diehard_2dsphere	8000	100	0.30209303	PASSED
diehard_3dsphere	4000	100	0.37641538	PASSED
diehard_squeeze	100000	100	0.00000000	FAILED
diehard_sums	100	100	0.19835615	PASSED
diehard_runs	100000	100	0.45844659	PASSED
diehard_runs	100000	100	0.44271558	PASSED
diehard_craps	200000	100	0.03758897	PASSED
diehard_craps	200000	100	0.07412468	PASSED
sts_monobit	100000	100	0.00000000	FAILED
sts_runs	100000	100	0.00000000	FAILED
sts_serial	100000	100	0.00678089	PASSED
sts_serial	100000	100	0.00000014	FAILED
rgb_minimum_distance	10000	1000	0.00000000	FAILED
rgb_permutations	100000	100	0.41729195	PASSED
rgb_lagged_sum	1000000	100	0.00000000	FAILED
rgb_kstest_test	10000	1000	0.00000060	FAILED
dab_bytedistrib	51200000	1	0.00000000	FAILED
dab_dct	50000	1	0.85055126	PASSED
dab_filltree	15000000	1	0.77145203	PASSED
dab_filltree	15000000	1	0.54588787	PASSED
dab_filltree2	5000000	1	0.82073965	PASSED
dab_filltree2	5000000	1	0.00000000	FAILED
dab_monobit2	65000000	1	1.00000000	FAILED

9.3.5 Blum Blum Shub

32 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.21687458	PASSED
diehard_operm5	1000000	100	0.00000000	FAILED
diehard_rank_32x32	40000	100	0.00000000	FAILED
diehard_rank_6x8	100000	100	0.00000000	FAILED
diehard_bitstream	2097152	100	0.00000000	FAILED
diehard_opso	2097152	100	0.00000000	FAILED
diehard_oqso	2097152	100	0.00000000	FAILED
diehard_dna	2097152	100	0.00000000	FAILED
diehard_count_1s_stream	256000	100	0.00000000	FAILED
diehard_count_1s_byte	256000	100	0.00000000	FAILED
diehard_parking_lot	12000	100	0.00000073	FAILED
diehard_2dsphere	8000	100	0.00096085	WEAK
diehard_3dsphere	4000	100	0.00283580	WEAK
diehard_squeeze	100000	100	0.00000000	FAILED
diehard_sums	100	100	0.00565991	PASSED
diehard_runs	100000	100	0.00000000	FAILED
diehard_runs	100000	100	0.00000000	FAILED
diehard_craps	200000	100	0.00000000	FAILED
diehard_craps	200000	100	0.00000000	FAILED
marsaglia_tsang_gcd	10000000	100	0.00000000	FAILED
marsaglia_tsang_gcd	10000000	100	0.00000000	FAILED
sts_monobit	100000	100	0.00000000	FAILED
sts_runs	100000	100	0.00000000	FAILED
sts_serial	100000	100	0.00000000	FAILED
sts_serial	100000	100	0.00000000	FAILED
rgb_minimum_distance	10000	1000	0.00000000	FAILED
rgb_permutations	100000	100	0.00000000	FAILED
rgb_lagged_sum	1000000	100	0.00000000	FAILED
rgb_kstest_test	10000	1000	0.00000000	FAILED
dab_bytedistrib	51200000	1	0.00000000	FAILED
dab_dct	50000	1	0.00000000	FAILED
dab_filltree	15000000	1	0.00000000	FAILED
dab_filltree	15000000	1	0.00000000	FAILED
dab_filltree2	5000000	1	0.00000000	FAILED
dab_filltree2	5000000	1	0.00000000	FAILED
dab_monobit2	65000000	1	1.00000000	FAILED

64 bit

Test Name	tsamples	psamples	p-value	Result
diehard_birthdays	100	100	0.19920411	PASSED
diehard_operm5	1000000	100	0.00000000	FAILED
diehard_rank_32x32	40000	100	0.00000000	FAILED
diehard_rank_6x8	100000	100	0.00000000	FAILED
diehard_bitstream	2097152	100	0.00000000	FAILED
diehard_opso	2097152	100	0.00000000	FAILED
diehard_oqso	2097152	100	0.00000000	FAILED
diehard_dna	2097152	100	0.00000000	FAILED
diehard_count_1s_stream	256000	100	0.00000000	FAILED
diehard_count_1s_byte	256000	100	0.00000000	FAILED
diehard_parking_lot	12000	100	0.00000073	FAILED
diehard_2dsphere	8000	100	0.00096085	WEAK
diehard_3dsphere	4000	100	0.17804038	PASSED
diehard_squeeze	100000	100	0.00000000	FAILED
diehard_sums	100	100	0.01043994	PASSED
diehard_runs	100000	100	0.00000000	FAILED
diehard_runs	100000	100	0.00000000	FAILED
diehard_craps	200000	100	0.00000000	FAILED
diehard_craps	200000	100	0.00000000	FAILED
marsaglia_tsang_gcd	10000000	100	0.00000000	FAILED
marsaglia_tsang_gcd	10000000	100	0.00000000	FAILED
sts_monobit	100000	100	0.00000000	FAILED
sts_runs	100000	100	0.00000000	FAILED
sts_serial	100000	100	0.00000000	FAILED
sts_serial	100000	100	0.00000000	FAILED
rgb_minimum_distance	10000	1000	0.00000000	FAILED
rgb_permutations	100000	100	0.00000000	FAILED
rgb_lagged_sum	1000000	100	0.00000000	FAILED
rgb_kstest_test	10000	1000	0.00000000	FAILED
dab_bytedistrib	51200000	1	0.00000000	FAILED
dab_dct	50000	1	0.00000000	FAILED
dab_filltree	15000000	1	0.00000000	FAILED
dab_filltree	15000000	1	0.00000000	FAILED
dab_filltree2	5000000	1	0.00000000	FAILED
dab_filltree2	5000000	1	0.00000000	FAILED
dab_monobit2	65000000	1	1.00000000	FAILED

10 Examples

10.1 Fitting data to a Cauchy distribution

On adding gaussian noise to 20 points sampled from a Cauchy with $x_0 = 23$, $\gamma = 7$, a fit is obtained with $x_0 = 24.0614$, $\gamma = 7.15664$




Figure 1: Fitting to a distribution

10.2 Estimating the value of π using Monte-Carlo methods

The code for the example can be found under *testing/test_montecarlo.cpp*




Figure 2: Monte-Carlo estimation of π

10.3 Path Tracer using Mersenne Twister to sample BRDFs




Figure 3: A Path-Tracer using DiceForge’s Mersenne Twister to sample light-ray directions

10.4 Chaos Game using LFSR to choose points




Figure 4: A Sierpinski Triangle formed through a Chaos Game using a regular triangle and a factor of $\frac{1}{2}$

10.5 Sampling a Maxwell-Boltzmann distribution

The code for the example can be found under *testing/test_maxwell.cpp*




Figure 5: PDF from sampled Maxwell-Boltzmann

11 Appendix

11.1 Parameters in Mersenne Twister

The optimum values for random number generators are :
32-bit

- $(w,n,m,r)=(32,624,397,31)$
- $a= 2573724191$
- $b= 0x9d2c5680$
- $c = 0xefc60000$
- $u = 11$
- $s = 7$
- $t = 15$
- $l = 18$

64-bit

- $(w,n,m,r)=(64,312,156,31)$
- $a= 0xB5026F5AA96619E9$
- $b= 0xD66B5EF5B4DA0000$
- $c = 0xFDED6BE000000000$
- $u = 29$
- $s = 17$
- $t = 37$
- $l = 41$

12 Sources

1. Random Numbers with LFSR (Linear Feedback Shift Register) - Computerphile (https://youtu.be/Ks1pw1X22y4?si=0Wj4uReK7nXS_0Bn)
2. XORShift RNGs by George Marsaglia (https://www.researchgate.net/publication/5142825_Xorshift_RNGs)
3. The Dieharder Tests by Robert G. Brown (<https://webhome.phy.duke.edu/~rgb/General/dieharder.php>)
4. Integration using Gaussian quadrature (<https://youtu.be/Hu6yqs0R7GA?si=mztKG7EjKaJvzI5W>)
5. Adaptive Gaussian quadrature (https://youtu.be/U4NUXAwR8E?si=0wpjGlyJlDe_V1S8)
6. Non-linear least squares regression (<https://www.uio.no/studier/emner/matnat/math/MAT3110/h19/undervisningsmateriale/lecture13.pdf>)
7. Cauchy distributions (<https://www.itl.nist.gov/div898/handbook/eda/section3/eda3663.htm>)

and of course Wikipedia